

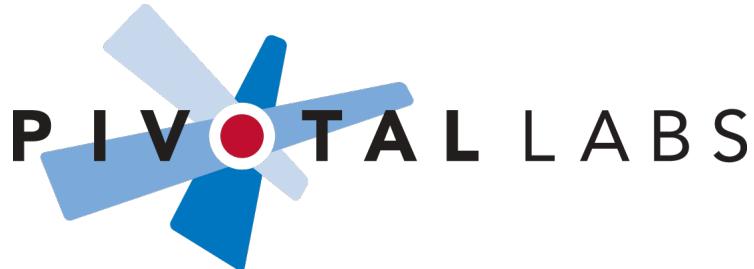
Getting Started with Cloud Foundry

Ian Huston, Data Scientist

Who am I?

- Ian Huston
- @ianhuston
- www.ianhuston.net
- Talk resources:
[https://github.com/ihuston/
python-cf-examples](https://github.com/ihuston/python-cf-examples)
- Data Scientist
- Use PyData stack for
predictive analytics and
machine learning
- Previously a theoretical
physicist using Python for
numerical simulations & HPC

Who are Pivotal?



Apache Tomcat



CLOUD
FOUNDRY™



OPEN DATA
PLATFORM



Pivotal
Big Data Suite



Pivotal™

Aims for this tutorial



Goals:

- Understand how to use CF as a developer/data scientist
- Understand how to push PyData apps and bind to services

Anti-goals:

- Understand CF from an operations viewpoint
- Install & run a full Cloud Foundry installation

Plan

- What is Cloud Foundry?
- Getting started with a hosted CF instance
- Pushing your first app
- What are buildpacks?
- Using PyData packages
- Using data services like Redis
- Putting it all together

Important Resources

Clone this repo:

<https://github.com/ihuston/python-cf-examples>

OR download and unpack this file:

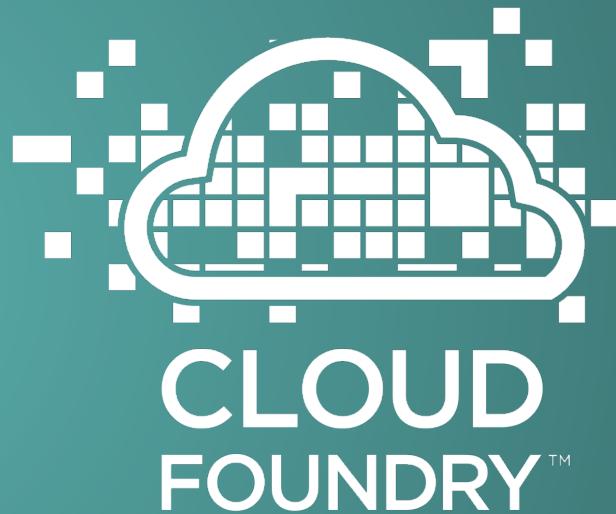
<http://tinyurl.com/cf-pydata>

What is Cloud Foundry?

<http://cloudfoundry.org>

Open Source
Multi-Cloud Platform

Simple App Deployment,
Scaling & Availability

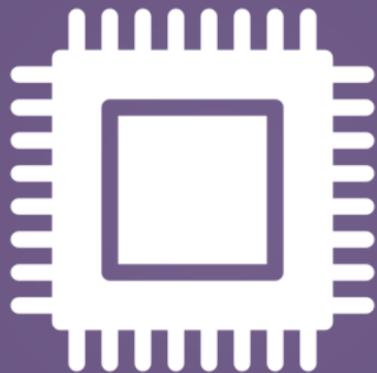


Cloud Applications Haiku

Here is my source code
Run it on the cloud for me
I do not care how.

- Onsi Fakhouri
[@onsijoe](https://twitter.com/onsijoe)

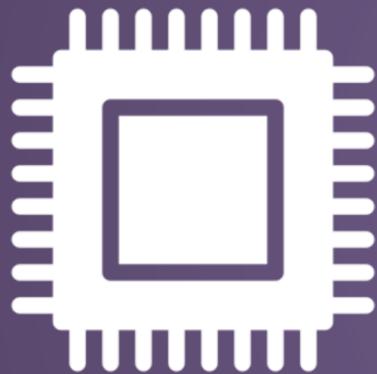
How can data scientists use CF?





Data Services

Easy control of incoming data



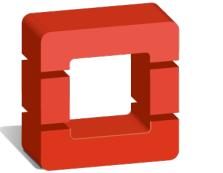
Distributed computation



Data Driven Applications

Multi-cloud

- Same applications running across different cloud providers:



openstack™



Private



Public



Pivotal
Web Services



IBM Bluemix™



CenturyLink®

Hosted
Pivotal

Hosted options



Pivotal
Web Services

- IBM Bluemix
- Anynines
- HP Helion Development Platform
- ActiveState Stackato
- CenturyLink AppFog

Getting Started on Pivotal Web Services

- Go to <http://run.pivotal.io>
- **Trial lasts for 60 days, no credit card required**
- Click ‘Sign Up’ and enter your email
- Open confirmation email and verify registration
- Complete SMS verification check
- Choose an organisation name (your name/handle is fine)

Downloading the Command Line Interface

- Register at <http://run.pivotal.io>
- In the CF management website click ‘Tools’ and download appropriate package
- Or go to <https://github.com/cloudfoundry/cli>
- Or use homebrew on OSX:

```
$ brew tap pivotal/tap
$ brew install cloudfoundry-cli
```
- Provides the cf command, your interface to Cloud Foundry

cf push

Challenge 1: Pushing your first app

Logging in to PWS

- Choose PWS API endpoint:

```
$ cf api https://api.run.pivotal.io
```

- Login:

```
$ cf login
```

- Enter your email address and password

- You should see output like:

```
API endpoint: https://api.run.pivotal.io (API version:2.28.0)
```

```
User:          YOUR_USER_NAME
```

```
Org:          YOUR_ORG_NAME
```

```
Space:        development
```

Challenge 1: Pushing your first application

- Choose PWS API endpoint: `$ cf api https://api.run.pivotal.io`
- Login: `$ cf login`
- Code directory: `01-simple-python-app`
- Deploy with

\$ cf push

- Check it worked at `http://myapp-RANDOM-WORDS.cfapps.io`
- Turn off your app when finished with `cf stop myapp` (but not yet!)

Simple Flask App Demo

- Simple one page “Hello World” web app
- Video: <https://www.youtube.com/watch?v=QOfD6tnoAB8>
- Demonstrates:
 - Installation of requirements
 - Scaling properties
- Need to Provide:
 - App files
 - Dependencies listed in requirements.txt file
 - Optional manifest.yml file with configuration for deployment

WHAT JUST HAPPENED?



\$ cf push

Source
Code

Browser

2. Set up domain

1. Upload code

Cloud
Controller

3. Install Python
&
Dependencies

C
F
R
O
U
T
E
R

Instance

Instance

5. Load balance
between
instances

5. Start app
and accept
connections

Send request to URL

4. Copy app into
containerised
instances

Pivotal™

What just happened?

1. Application code is uploaded to CF
 2. Domain URL is set up ready for routing
 3. Cloud controller builds application in container:
 - Python interpreter selected
 - Dependencies installed with pip
 4. Container is replicated to provide instances
 5. App starts and Router load balances requests
- See what's happening using logs: `$ cf logs myapp --recent`

Python on Cloud Foundry

- First class language (with Go, Java, Ruby, Node.js, PHP)
- Automatic app type detection
 - Looks for requirements.txt or setup.py
- **Buildpack** takes care of
 - Detecting that a Python app is being pushed
 - Installing Python interpreter
 - Installing packages in requirements.txt using pip
 - Starting web app as requested (e.g. python myapp.py)

Scaling memory and instances

- We can scale our application as needed in terms of memory and number of instances:

```
$ cf scale myapp -i 5
```

```
$ cf scale myapp -m 256M
```

- Check app in browser to see different ports being used.
- Scale back down with \$ cf scale myapp -i 1

How does this work?

- Containerised application is cached and ready to be deployed.
- Scaling number of instances replicates container and load balances requests across all instances.
- Scaling memory requires restarting app.
- Auto-scaling is also possible.

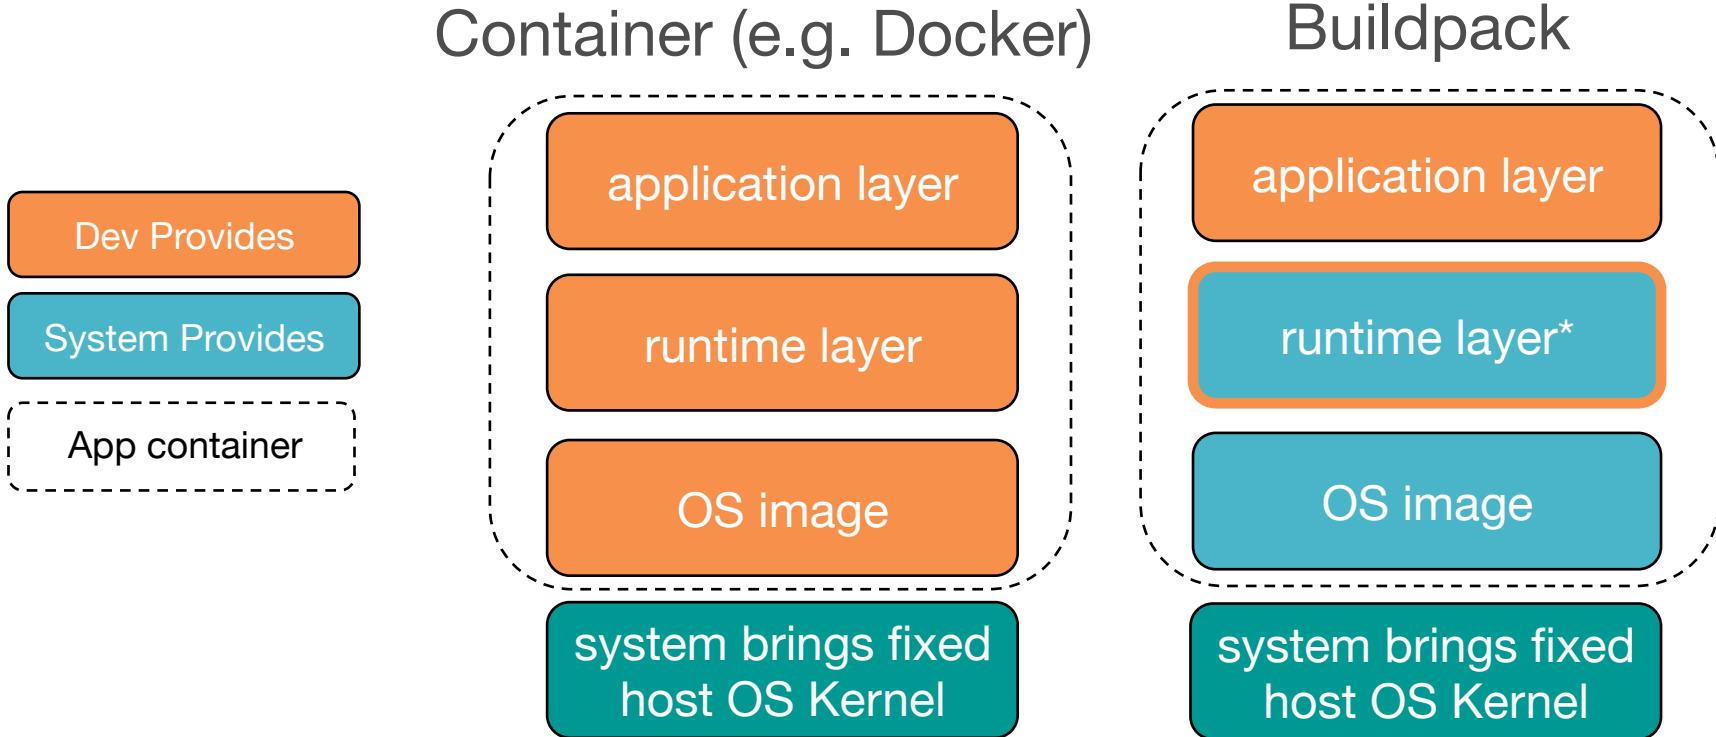
Buildpacks

Challenge 2: Pushing PyData apps

What are buildpacks?

- Idea and format from Heroku
- Responsible for doing whatever is necessary to get your app running.
- Buildpacks take care of
 - Detecting which type of application is being pushed
 - Installing the appropriate run-time
 - Installing required dependencies or other artifacts
 - Starting the application as requested
- Official buildpacks for Python, Java, Node.js, Go, Ruby, PHP & for static websites and running binaries

Containers vs Buildpacks



* Devs may bring a custom buildpack

Custom buildpacks

- Instead of using an official buildpack you can use any custom buildpack installed on your CF or available on Github.
- Only 3 shell scripts needed:
 - detect
 - compile
 - release
- Specify buildpack with -b or in manifest.yml

Community Buildpacks

[https://github.com/cloudfoundry-community/
cf-docs-contrib/wiki/Buildpacks](https://github.com/cloudfoundry-community/cf-docs-contrib/wiki/Buildpacks)

- Lots of languages:
Clojure, Haskell, .NET, Erlang, Elixir, etc.
- RShiny app buildpack:
<https://github.com/alexkago/cf-buildpack-r>
- Can also use some Heroku buildpacks without modification.

Official Python buildpack

- ✓ Great for simple pip based requirements
- ✓ Well tested and officially maintained
- ✓ Covers both Python 2 and 3
- ✗ Suffers from the Python Packaging Problem:
 - Hard to build packages with C, C++ or Fortran extensions
 - Complicated local configuration of libraries and paths needed
 - Takes a long time to build main PyData packages from source

Using conda for package management

- <http://conda.pydata.org>
- Benefits:

- Uses precompiled binary packages
- No fiddling with Fortran or C compilers and library paths
- Known good combinations of main package versions
- Really simple environment management (better than virtualenv)
- Easy to run Python 2 and 3 side-by-side



CONTINUUM
ANALYTICS

Go try it out if you haven't already!

How to use the conda buildpack

<https://github.com/ihuston/python-conda-buildpack>

- Specify as a custom buildpack when pushing app with manifest or -b command line option.
- Export your current environment to a environment.yml file
- Or write requirements.txt (pip) and conda_requirements.txt
- Send me feedback & pull requests!

Challenge 2: Pushing a PyData app

- Code directory: 02-pydata-spyre-app
- Spyre – Adam Hajari <https://github.com/adamhajari/spyre>
- This app uses the Pydata buildpack to install Matplotlib, NumPy and more.
- Spyre provides a simple way to build interactive web based visualisations similar to Rshiny.

Using Services

Challenge 3: Using Redis in an app

Services

These are available on Pivotal CF (Pivotal's packaged Cloud Foundry offering).



See <http://run.pivotal.io> for the services available on Pivotal Web Services.

Cloud Native Applications

- Suitable for deployment on cloud platforms
- Can scale up easily without changes
- Follow these rules: <http://12factor.net>
- In a nutshell:
 - Apps are stateless processes
 - Easy to create and destroy apps with no side-effects
 - Persistent state handled through backing services
 - Interact with services through port binding

Challenge 3: Binding a service to an app

- Create a free Redis service:

```
$ cf cs rediscloud 30mb myredis
```

- Bind to our earlier app:

```
$ cf bind-service myapp myredis
```

- See how apps find credentials in environmental variables:

```
$ cf env myapp
```

Challenge 3: Using Services in an app

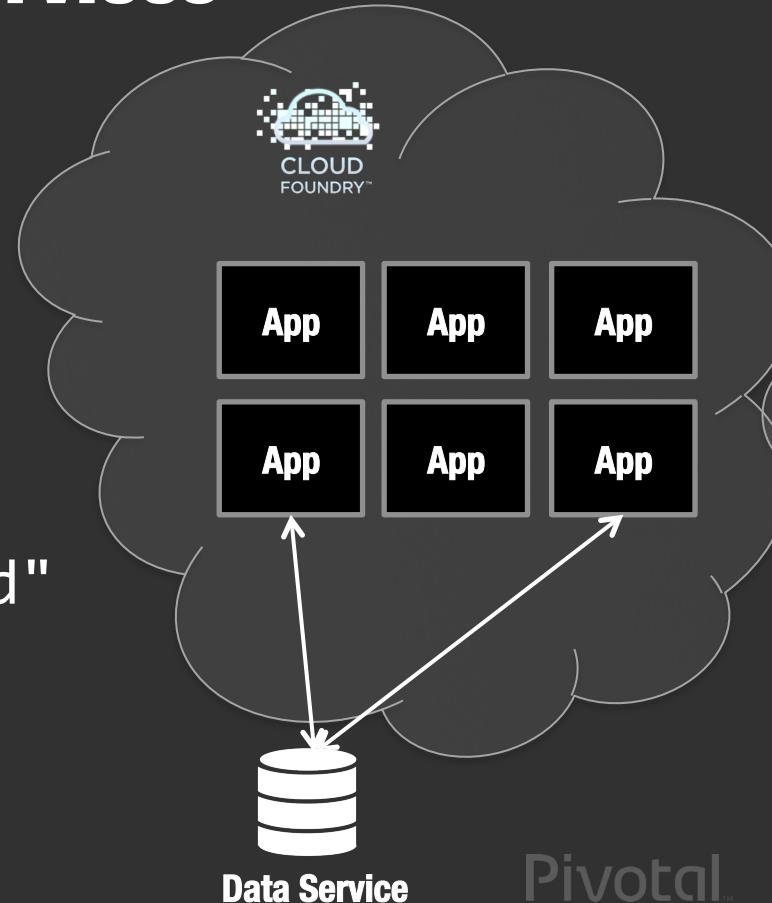
- Code directory: 03-services-redis
- This app binds to a Redis service if available and counts the number of hits on the app homepage.
- Add keys and values by adding /keyname/value to URL.
- Data is persisted in Redis, will survive deleting and restarting the app.
- Multiple instances all access the same service.

User Provided Services

How to add User Provided Services:

Standalone Hadoop or Apache Spark cluster,
Big Data System, RDBMS etc.

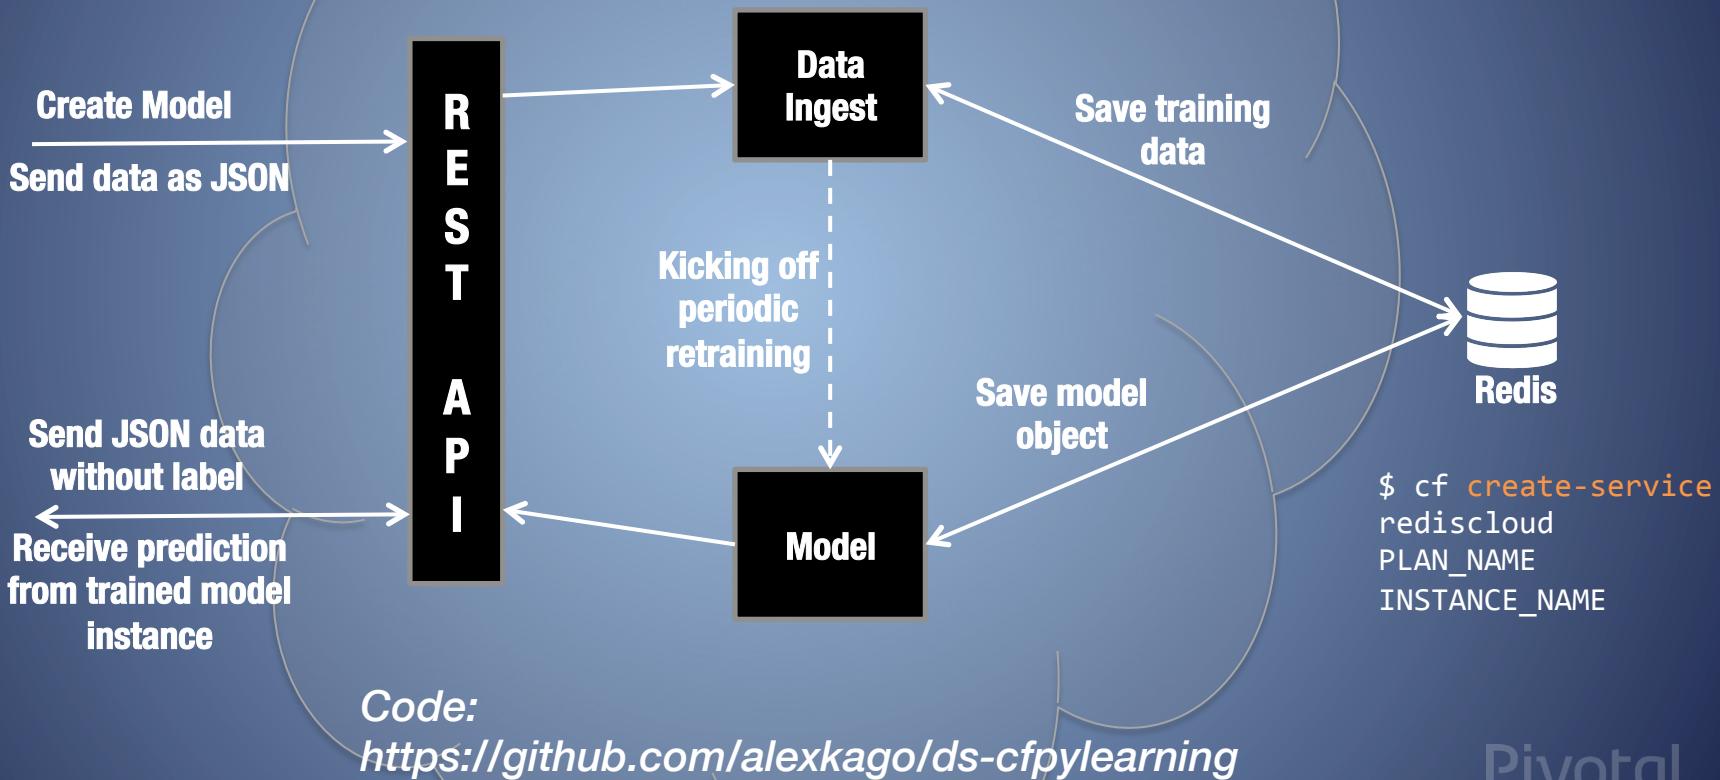
```
$ cf cups SERVICE_INSTANCE -p  
"host, port, username, password"
```



Putting it all together

Challenge 4: Build your own
prediction API

PREDICTION API ARCHITECTURE



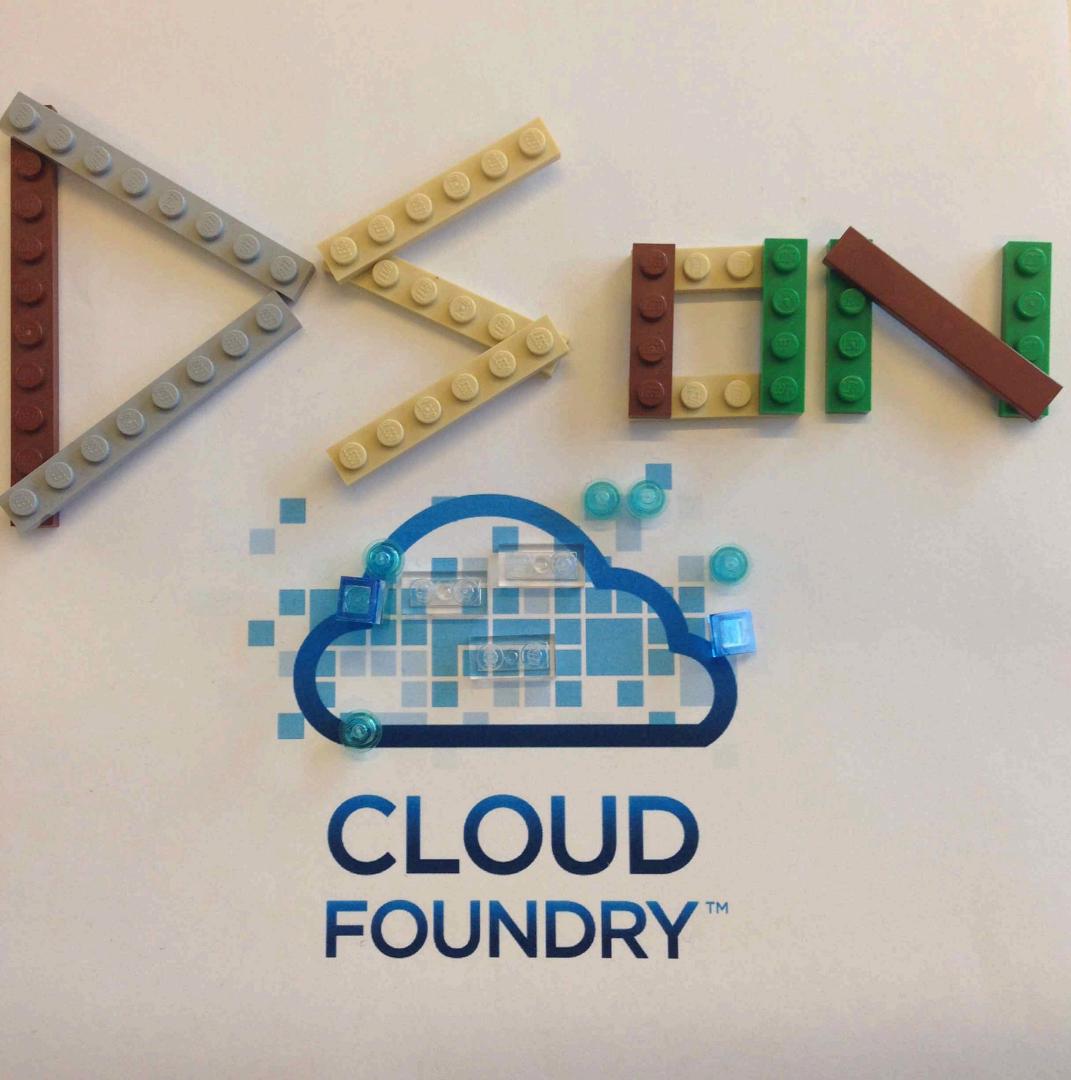
Pivotal™

Final challenge: Build your own predictive API

- Code directory: 04-learning-api
- Simple Flask + scikit-learn based machine learning API
- Push the application and go to the app URL to see instructions on how to use.
- Model is built in scikit-learn and persisted in Redis
- Simplified version of this project by Alex Kagoshima:
<https://github.com/alexkago/ds-cfpylearning>

Resources

- Docs: <http://docs.cloudfoundry.org>
- CF Summit videos: <http://cfsummit.com>
- Join the CF community: <http://cloudfoundry.org>
- CF meetups: <http://cloud-foundry.meetup.com>
- Don't forget to stop your apps with `cf stop myapp.`



Show off your data science related Cloud Foundry apps:

Twitter: @dsoncf
<http://dsoncf.com>

@ianhuston

